

Tokenization and Embedding in Large Language Models (LLMs)

Dr. Saikat Sarkar

Assistant Professor

Bangabasi College, University of Calcutta, Kolkata, India

to.saikatsarkar17@gmail.com

Contents

- LLMs
- LLM workflow
- How LLMs Understand Text
- Tokenization
- Embedding
- Positional Embedding

What is an LLM (Large Language Model)?

- A Large Language Model (LLM) is an advanced **AI model**
- Trained on **vast amounts of text** (books, websites, and other sources)
- To understand, generate, and reason with **human language**.

LLM: Applications (Typical)

- **Question answering**
 - – Example: You can ask, “What is the capital of India?” and the model will reply “New Delhi”
- **Writing essays, code**
- **Translating languages**
- **Generating creative text**
 - – Example: The model can write a poem, story, or dialogue between two characters.


Big players

- ChatGPT: by OpenAI
- Big tech
 - Gemini: Google's version
 - Meta AI: Meta's version
 - Copilot: Microsoft's version
- Startups ---
 - Claude: Anthropic's version
 - Grok: xAI's version
 - Perplexity
- DeepSeek (Chinese Co.)
- Le Chat: Mistral's version (French Co.)

LLM Workflow



LLM Workflow




Criteria for Interpretation in LLMs

- **Task Type** Recognition:
 - Write an email
 - Write a story
- **Intent** Detection:
 - Explain photosynthesis: Command
 - What is photosynthesis? Question
- **Polarity** and **Sentiment** Recognition:
 - This movie was amazing: positive
 - This movie was terrible: negative
- **Context:**
 - He sat on the bank: river context
 - He went to the bank: finance context

How LLMs Understand Text


- LLMs **can't understand** words or sentences directly — they work with **numbers**.
- Before feeding text to the model, we need to **convert** text into numbers that the model can process.
- That's where **Tokenization** and **Embedding** come in!
 - **Tokenization** – Split text into tokens.
 - **Embedding** – Convert tokens into numerical vectors.

LLM Workflow: Tokenization and Embedding



What is Tokenization?

- **Tokenization** means breaking a sentence or text into smaller parts called **tokens**.
- These tokens can be:
 - Words
 - Subwords (parts of words)
 - Characters
 - Punctuation marks



Tokenization: Examples

- LLM vs llm
- LLM!
- $127 + 2895 * 164 = 474907$
- $11+2*10=31$
- $111+2*100=311$

```
for i in range (10):
    print(i)
    if i> 5:
        exit(0)
```

Tokenization: Character Level

- Why not just use each character as a token?
 - “I am learning LLM”: 4 words, 17 characters
- LLMs process sequences token by token.
 - More tokens = longer sequences = more memory + compute time.
- Loss of semantic information
 - “unhappiness” → characters [u, n, h, a, p, p, i, n, e, s, s]
 - Each character gives almost no clue about meaning.

Tokenization: Word Level

- Why not just split every sentence by spaces and assign each word an ID?
- **Vocabulary explosion** (too many unique words)
 - “play”, “played”, “playing” → unrelated in word-level vocab.
- A word-level model cannot handle **unseen** words.
 - If the training data never saw “cyber-bioinformatics”, the model can’t represent it.
 - It will either replace it with [UNK] (unknown)
 - Or fail to interpret.

Modern LLMs use **subword tokenization**

Word	Sub-word tokens
play	["play"]
playing	["play", "ing"]
replay	["re", "play"]
replaying	["re", "play", "ing"]

Byte Pair Encoding (BPE)

- Suppose the data to be encoded is “aabdaaaabac”
- Replace the most occurred pair “aa” with Z

ZabdZabac
Z=aa

- Repeat: replace “ab” with Y

ZYdZYac
Y=ab
Z=aa

- Repeat: replace “ZY” with X

XdXac
X=ZY
Y=ab
Z=aa

WordPiece Tokenization

- Similar to the BPE
- Now, instead of just choosing the most frequent pair (like BPE), WordPiece uses a **likelihood-based criterion**.
- For each candidate pair (x, y) , WordPiece estimates:

$$\text{score}(x, y) = \frac{\text{freq}(xy)}{\text{freq}(x) \times \text{freq}(y)}$$

- Merge the best pair based on the score

SentencePiece Tokenizer


- BPE and WordPiece assume spaces = word boundaries
- That works well for English, but fails for:
 - Languages **without spaces** (e.g., Chinese, Japanese, Thai),
- Treats the **entire text** as a sequence of **Unicode characters**
- Keeps or replaces spaces with a special symbol, typically "**_**"
“I love learning” -> “_I_learn_love_learning”
Could merge like “_love” + “_learn” → “_love_learn” (possible if common in corpus)

Tokenization: Non-textual Data

- Image: Sequence of patches



- Audio:




- Video

Questions?

From Tokens to Embeddings

Models **don't use** IDs directly — they use embeddings.

Embedding: a numerical vector that represents the meaning of a token.



Embedding Matrix:


[

[0.19, -0.41, 0.88, ...],
[0.73, -0.05, -0.32, ...],
[-0.11, 0.25, 0.66, ...],

.....

]

Visualizing Embedding Space



Word	Vector embedding (512 dimensions)
cat	[1.5, -0.4, 7.2, 19.6, 20.2, ...]
dog	[1.7, -0.3, 6.9, 19.1, 21.1, ...]
fish	[-5.2, 3.1, 0.2, 8.1, 3.5, ...]
fisherman	[-4.9, 3.6, 0.9, 7.8, 3.6, ...]
triangle	[60.1, -60.3, 10, -12.3, 9.2, ...]
PS4 Remote	[81.6, -72.1, 16, -20.2, 102, ...]

Learning Embeddings

1. Start with random embeddings

2. Model predicts next word

“The cat sat on the ___”

3. Compute error (loss)

4. Backpropagate and update embeddings

5. Similar words → similar embeddings

Token	Initial Embedding (simplified)
“cat”	[0.2, -0.4, 0.1]
“dog”	[-0.3, 0.5, -0.1]
“mat”	[0.8, -0.2, 0.3]

Positional Embedding

- Transformers process all tokens **in parallel**, not one after another.
- The model doesn't know the **order** of the words!
 - The **cat** chased the **dog**
 - The **dog** chased the **cat**
- We have to **tell the model** *where each word appears in the sequence*.
- A **positional embedding** is an **additional vector** added to each token's embedding that **represents its position** (index) in the sentence.


Positional Embedding

- Final input to the LLM = Token Embedding + Positional Embedding

Token	Token Embedding	Position	Positional Embedding	Final Input
“The”	[0.2, -0.1, 0.5]	1	[0.01, 0.99, 0.05]	[0.21, 0.89, 0.55]
“cat”	[0.8, -0.2, 0.3]	2	[0.02, 0.95, 0.10]	[0.82, 0.75, 0.40]
“sat”	[0.4, -0.4, 0.7]	3	[0.03, 0.90, 0.15]	[0.43, 0.50, 0.85]

Questions?

LLM Workflow: Tokenization and Embedding



Funny Things

- Counting Letters or Characters
 - How many letters are in the word ‘banana’?”
- Counting Words Accurately
 - “Repeat the word ‘yes’ 50 times.”
 - The model predicts tokens, not actual word counts

References

- Intro to Large Language Models by Andrej Karpathy (https://www.youtube.com/watch?v=zjkBMFhNj_g)
- Introduction to large language models by Google Cloud Tech (<https://www.youtube.com/watch?v=zizonToFXDs>)
- Transformers, explained: Understand the model behind ChatGPT By [Leon Petrou](https://www.youtube.com/watch?v=Pnd8bCJ4Z3A) (<https://www.youtube.com/watch?v=Pnd8bCJ4Z3A>)
- Hands-On Large Language Models by Jay Alammar, Maarten Grootendorst
- Build a Large Language Model (From Scratch) by Sebastian Raschka

Thanks!!!